

UKRAINIAN CATHOLIC UNIVERSITY

FACULTY OF APPLIED SCIENCES

DATA SCIENCE MASTER PROGRAMME

Visualization of high-dimensional data using t -SNE algorithm

Linear Algebra final project report

Authors:

Serhii BRODIUK

Andrii YURKIV

24 January 2019



APPLIED
SCIENCES
FACULTY ●

Abstract

In this report we describe our implementation of t -SNE algorithm for visualization of multidimensional data.

1 Introduction

Data visualization is one of the most important parts of applied data analysis nowadays. Without proper visualization it's hard, or sometimes even impossible to interpret the data in the appropriate way or make reasonable inferences regarding its structure. Dealing with data sets that contain more than 3 attributes start causing problems, because standard visualization techniques does not apply to the data in 4-dimensional space. For humans, it's even hard to perceive graphical information in 3-dimensional space, that's why the majority of classical techniques of data visualization are restricted to two dimensions.

It should be noted that one of the main purposes of visualization of high-dimensional data sets is to determine clusters inside the data which may be an important basement for further data analysis.

2 Motivation

Modern data often contains hundreds or even thousands of features. And classical visualization approaches cannot be used to determine patterns or clusters in it. This is the reason why so many different techniques for visualization high-dimensional data have been proposed. Majority of these techniques merely provide tools to display multidimensional data in two or three-dimensional form and do not give any inferences regarding the inner structure of the data. This severely limits the applicability of those techniques to real-world data, which sometimes contain too much features to reduce them nicely.

Another approach to this problem lies in the field of dimensionality reduction. We can apply classical plotting methods (for example, scatter plots) for the data obtained by decreasing the dimension of the original data set. Dimensionality reduction techniques differ substantially from visualization techniques because they are aimed not to make data more visually appealing or understandable, but to preserve as much of the significant structure of the high-dimensional data as possible in its low-dimensional representation [1].

3 Problem setting

The most important goal of visualization of multidimensional data is to correctly identify clusters of similar data points. This is crucial when we are working on unsupervised learning problem and need to know for sure what number of distinct clusters can our data set be divided in without losing valuable information about inner structure of our data.

As we will see in this report, many classical visualization algorithms correctly identify the relationships between data points and position similar ones together. But at the same time from visualizations those algorithms provide it's hard to tell where one class of points ends and other begins (boundaries between classes are not well-defined). And while this is not a problem for small data sets, for large ones it can cause serious obstacles for correct definition of of the number and location of data clusters.

t-SNE (*t*-distributed Stochastic Neighbor Embedding) algorithm is aimed to overcome those limitations by identifying the relevant patterns using the approach in which similar objects are modeled by nearby points and dissimilar objects are modeled by distant points with high probability [2].

4 Related work

5 Approach to Solution

6 Solution

7 Evaluation

Since *t*-SNE is primarily visualization technique, in order to evaluate its performance and visualizational capacity we need to compare visualization produced by *t*-SNE with the ones obtained using other non-parametric visualization techniques for multidimensional data, such as Multidimensional scaling using SMACOF (Scaling by Majorizing a Complicated Function) algorithm, Isomap, Locally Linear Embedding, Principal Components Analysis and Laplacian Eigenmaps. PCA is not usually used for visualization, but we decided that it would be a good idea to show how the visualizations obtained from manifold learning algorithms differ from those obtained from classic dimensionality reduction technique.

7.1 Experiments

In our work we experimented with 2 different data sets: Olivetti faces data set and handwritten digits data set.

Olivetti faces dataset contains 400 64×64 grayscale images of faces (10 images for each of 40 subjects). Those images were taken at different time with different details varying, such as face expression, lighting and facial details (absence or presence of glasses). They are labeled according to the identity of the person depicted.

Handwritten digits data set consists of 1797 8×8 grayscale images of handwritten digits of 10 classes. There are nearly 180 instances per digit class in this data set.

Technique	Parameters
<i>t</i> -SNE	<i>perplexity</i> = 5, 1000 iterations
MDS (SMACOF)	-
Isomap	<i>neighbors</i> = 5
LLE	<i>neighbors</i> = 12
PCA	-
Laplacian Eigenmaps	<i>neighbors</i> = 12

Table 1: Cost function parameters for the experiments

7.2 Experimental setup

In the original *t*-SNE paper [1], the authors use PCA to reduce dimensionality before feeding the data into algorithms that transform it into two-dimensional representation. They are doing it for the purpose of computation efficiency and it absolutely makes sense when we apply those algorithms to large data sets. But since the data sets which we use in our experiments are relatively small, preliminar dimensionality reduction doesn’t make much difference in terms of computational costs in our case, so we decided to pass this step up.

For each of these data sets, there is information about the class of each data point. Class information is only used for visualization purposes and is not considered during calculation of the spatial coordinates of the points in two-dimensional space. We decided to visualize the whole images instead of points on scatter plots because this kind of plot is much more understandable and provides better means of evaluating how well the mapping preserves the similarities within each class and how accurately it separates one class from the other.

We used library utilities in order to demonstrate visualization capacity and performance of algorithms, with which we are comparing our implementation of *t*-SNE algorithm. We used `Isomap`, `LocallyLinearEmbedding`, `MDS` and `SpectralEmbedding` classes from package `sklearn.manifold` and `PCA` class from package `sklearn.decomposition`.

The cost function parameters we used in our experiments are listed in Table 1. Here *perplexity* represents the perplexity of the conditional probability distribution induced by the Gaussian kernel and *neighbors* represents the number of nearest neighbors employed in a neighborhood graph. Multidimensional scaling was performed using SMACOF algorithm with 4 initializations and 300 iterations and without any other explicit parameters. PCA uses regular library settings.


For each algorithm we visualized all data points present in the data set, that’s why some of them look a bit crowded with too many images.

7.3 Results


In figures 1-3 we displayed the visualization produced by t -SNE, MDS, Isomap, LLE, PCA and Laplacian Eigenmaps algorithms applied to the Olivetti faces data set. Those result support our intuition that t -SNE strongly outperforms other techniques in visualization of multidimensional data set. For example, MDS generates ball-like figure, in which only few face classes are visually separated from the other images. Almost all algorithms have few face classes that are correctly clustered and located so that their group is distinct from the others. But at the same they have the majority of images overlapping, which makes it hard to distinguish the borders of the clusters. With such extreme overlaps as in plots produced by LLE or Laplacian Eigenmaps those graphs are almost useless, because if we hadn't class information we would fail to both number of distinct classes and their location.

On the contrary, t -SNE creates a map with relatively accurate separation of clusters. It locates images of the same person very close to each other, even if some distractions are introduced (like glasses or face rotation). We can observe a number of sole points that don't belong to any cluster. Those are the faces with some specific facial expressions, which t -SNE is not able to capture. Besides it, there are some misclassified data points, but most of them correspond to distorted images which is hard to identify.

8 Conclusions




(a) using t -SNE




(b) using MDS

Figure 1: Visualization of Olivetti faces data set




(a) using Isomap




(b) using LLE

Figure 2: Visualization of Olivetti faces data set




(a) using PCA




(b) using Laplacian Eigenmaps


Figure 3: Visualization of Olivetti faces data set




(a) Visualization by t -SNE




(b) Visualization by MDS




(c) Visualization by Isomap



(d) Visualization by LLE



(e) Visualization by PCA



(f) Visualization by LE

Figure 4: Visualizations of handwritten digits data set